

# EECE 5698 Assignment3: Parallel Classification

Zifeng Wang

## Problem 1

(1a) `SparseVector` is a subclass of standard python dictionary, which means it shares the same methods with the standard dictionary. Also, it implements new methods so that it extends the functionalities of standard dictionary.

(1b) Methods it shares with standard dictionary : has `key()`, `items()`, `keys()`, etc.  
Methods it has that are not present in a dictionary: `safeAccess(key)`, `dot(other)`,  
`__add__(other)`, `__sub__(other)`, `__mul__(s)`, `__rmul__(s)`

(1c) The operations and outcomes are as below

```
In [1]: from SparseVector import SparseVector as SV
In [2]: sv1 = SV([(1,3), (5,8)])
In [3]: sv2 = SV([(1,10), (3,6)])
In [4]: sv1.dot(sv2)
Out[4]: 30
In [5]: sv2.dot(sv1)
Out[5]: 30
In [6]: 5 * sv1
Out[6]: {1: 15, 5: 40}
In [7]: sv1 * 5
Out[7]: {1: 15, 5: 40}
In [8]: 5 * sv2
Out[8]: {1: 50, 3: 30}
In [9]: sv2 * 5
Out[9]: {1: 50, 3: 30}
```

## Problem 2 Here we prove 2(b) first

2(b)

For simplicity, prove 2(b) first

$$\begin{aligned}\nabla \ell(\beta; x, y) &= \frac{1}{1 + e^{-y\beta^T x}} e^{-y\beta^T x} (-yx) \\ &= -\frac{e^{-y\beta^T x} e^{y\beta^T x}}{(1 + e^{-y\beta^T x}) e^{y\beta^T x}} (yx) \\ &= -\frac{yx}{1 + e^{y\beta^T x}}\end{aligned}$$

2(a)

Then we prove 2(a), we could write the gradient in the form of components of  $\beta$  and  $x$ , where  $\beta = [\beta_1, \dots, \beta_d]^T$ ,  $x = [x_1, \dots, x_d]^T$ . Let  $i \in 1, \dots, d$ .

$$\frac{\partial \ell(\beta_i; x, y)}{\partial \beta_i} = -\frac{yx_i}{(1 + e^{y\beta^T x})^2}$$

So we have

$$\begin{aligned}\frac{\partial^2 \ell(\beta_i; x, y)}{\partial \beta_i \partial \beta_j} &= -\frac{yx_i}{(1 + e^{y\beta^T x})^2} (-e^{y\beta^T x}) yx_j \\ &= \frac{y^2 e^{y\beta^T x}}{(1 + e^{y\beta^T x})^2} x_i x_j \\ &= k x_i x_j\end{aligned}$$

Where  $k = \frac{y^2 e^{y\beta^T x}}{(1 + e^{y\beta^T x})^2} > 0$  does not depend on  $i, j$

So we have

$$\nabla^2 \ell(\beta; x, y) = kxx^T$$

Which is positive semi definite, because for  $m \in \mathbb{R}^d$

$$m^T (kxx^T) m = k(mx^T)^2 \geq 0$$

So we have the loss is convex. From the definition of convex function

2(c)

From 2(a) we have

$$\frac{\partial \ell(\beta_j; x, y)}{\partial \beta_j} \Big|_{x_j=0} = -\frac{0y}{(1 + e^{y\beta^T x})^2} = 0$$

Which means  $\nabla \ell(\beta; x, y)$  does not rely on  $\beta_j$  given  $x_j = 0$

Similarly we have

$$\begin{aligned}\ell(\beta; x, y) &= (1 + e^{-y\beta^T x}) \\ &= (1 + e^{-y \sum_{i=1}^d x_i \beta_i}) \\ &= (1 + e^{-y \sum_{i=1, i \neq j}^d x_i \beta_i} e^{-y x_j \beta_j})\end{aligned}$$

When  $x_j = 0$ ,  $e^{-y x_j \beta_j} = 1$ ,  $\ell(\beta; x, y)$  does not rely on  $\beta_j$  as well.

2(d)

From the conclusions of previous problems, we get

$$\sum_{i=1}^n \ell(\beta; x_i, y_i)$$

is convex as it is just sum of convex functions  $\ell(\beta; x_i, y_i)$ .

Also, we have  $\lambda \geq 0$  and we know all norms are convex, we have

$$\nabla^2(\lambda \|\beta\|_2^2) = 2\lambda I$$

So  $\lambda \|\beta\|_2^2$  is convex as well. As a result we have

$$L(\beta) = \sum_{i=1}^n \ell(\beta; x_i, y_i) + \lambda \|\beta\|_2^2$$

is convex.

### Problem 3

3(a) Here below are the function definitions of `logisticLoss`, `gradLogisticLoss` and `gradTotalLoss`

```
def logisticLoss(beta,x,y):
    """
        Given sparse vector beta, a sparse vector x, and a binary value y in {-1,+1}, compute the logistic loss
        l(B;x,y) = log( 1.0 + exp(-y * <β,x>) )

    The input is:
        - beta: a sparse vector β
        - x: a sparse vector x
        - y: a binary value in {-1,+1}

    """
    return np.log(1.0 + np.exp(-y * x.dot(beta)))

def gradLogisticLoss(beta,x,y):
    """
        Given a sparse vector beta, a sparse vector x, and
        a binary value y in {-1,+1}, compute the gradient of the logistic loss
        ∇l(B;x,y) = -y / (1.0 + exp(y <β,x>)) * x

    The input is:
        - beta: a sparse vector β
        - x: a sparse vector x
        - y: a binary value in {-1,+1}

    """
    return -y / (1.0 + np.exp(y * x.dot(beta))) * x

def gradTotalLoss(data,beta, lam = 0.0):
    """
        Given a sparse vector beta and a dataset perform compute the gradient of regularized total logistic loss :
        ∇L(β) = Σ_{(x,y) in data} ∇l(β;x,y) + 2λ β

    Inputs are:
        - data: a python list containing pairs of the form (x,y), where x is a sparse vector and y is a binary value
        - beta: a sparse vector β
        - lam: the regularization parameter λ
    """

    gradLoss = SparseVector({})
    for (x, y) in data:
        gradLoss += gradLogisticLoss(beta, x, y)
    return gradLoss + 2.0*lam*beta
```

### 3(b) Here below is the final function definition of test

```

def test(data,beta):
    """ Output the quantities necessary to compute the accuracy, precision, and recall of the prediction of labels in a dataset under a given  $\beta$ .
    The accuracy (ACC), precision (PRE), and recall (REC) are defined in terms of the following sets:
        P = datapoints (x,y) in data for which  $y > 0$ 
        N = datapoints (x,y) in data for which  $y \leq 0$ 
        TP = datapoints in (x,y) in P for which  $y=1$ 
        FP = datapoints in (x,y) in P for which  $y=-1$ 
        TN = datapoints in (x,y) in N for which  $y=-1$ 
        FN = datapoints in (x,y) in N for which  $y=+1$ 

    For #XXX the number of elements in set XXX, the accuracy, precision, and recall of parameter vector  $\beta$  over data are defined as:
        ACC( $\beta$ ,data) = (#TP+#TN) / (#P + #N)
        PRE( $\beta$ ,data) = #TP / (#TP + #FP)
        REC( $\beta$ ,data) = #TP / (#TP + #FN)

    Inputs are:
        - data: an RDD containing pairs of the form (x,y)
        - beta: vector  $\beta$ 

    The return values are
        - ACC, PRE, REC
    .....
    P = 0
    N = 0
    TP = 0
    FP = 0
    TN = 0
    FN = 0
    for x, y in data:
        temp = x.dot(beta)
        if temp > 0:
            P += 1
            if y == 1:
                TP += 1
            elif y == -1:
                FP += 1
        else:
            N += 1
            if y == -1:
                TN += 1
            elif y == 1:
                FN += 1

    ACC = (TP + TN) / float(P + N)
    PRE = TP / float(TP + FP)
    REC = TP / float(TP + FN)
    return ACC, PRE, REC

```

### 3(c)

We choose  $\lambda$  from {0.0, 1.0, 10.0}

$$\lambda = 0.0$$

```

(Yuang.Zhang@ccf0013 Assignment5)$ python LogisticRegression.py mushrooms/mushrooms_train --testdata mushrooms/mushrooms_test --lam 0.0 --max_iter 20
Reading training data from mushrooms/mushrooms_train
Read 7416 data points with 117 features in total
Reading test data from mushrooms/mushrooms_test
Read 1000 data points with 115 features
Training on data from mushrooms/mushrooms.train with  $\lambda = 0.0$ ,  $\epsilon = 0.1$  , max iter = 20
k = 0 t = 3.354545294189 L( $\beta$ , $\lambda$ ) = 5140.379451033264 ||V(L( $\beta$ , $\lambda$ ))||_2 = 4273.0823026041 gamma = 0.000470184284576 ACC = 0.390 PRE = 0.853235914749 REC = 1.0
k = 1 t = 1.153560686686 L( $\beta$ , $\lambda$ ) = 2516.999544902876 ||V(L( $\beta$ , $\lambda$ ))||_2 = 1072.956450920276 gamma = 0.000470184284477 ACC = 0.32 PRE = 0.892123247671 REC = 0.968491486989
k = 2 t = 1.03371118912 L( $\beta$ , $\lambda$ ) = 1072.956450920276 ||V(L( $\beta$ , $\lambda$ ))||_2 = 802.3372848311758 gamma = 0.00362797066 ACC = 0.979 PRE = 0.965765757576 REC = 0.996282577881
k = 3 t = 13.5136380196 L( $\beta$ , $\lambda$ ) = 772.3391480787859 ||V(L( $\beta$ , $\lambda$ ))||_2 = 980.5492701241932 gamma = 0.000470184284576 ACC = 0.974 PRE = 0.981293007519 REC = 0.970360223948
k = 4 t = 16.9038770199 L( $\beta$ , $\lambda$ ) = 648.96895293955 ||V(L( $\beta$ , $\lambda$ ))||_2 = 636.2301783321042 gamma = 0.000470184284576 ACC = 0.976 PRE = 0.972426470588 REC = 0.983271375465
k = 5 t = 20.1671419144 L( $\beta$ , $\lambda$ ) = 583.1293093296519 ||V(L( $\beta$ , $\lambda$ ))||_2 = 390.12150441871074 gamma = 0.00078364164096 ACC = 0.975 PRE = 0.981238273921 REC = 0.972138959108
k = 6 t = 23.5557079315 L( $\beta$ , $\lambda$ ) = 552.8636871007043 ||V(L( $\beta$ , $\lambda$ ))||_2 = 474.3895736703113 gamma = 0.000470184284576 ACC = 0.981 PRE = 0.97790552486 REC = 0.98698847584
k = 7 t = 26.8180449006 L( $\beta$ , $\lambda$ ) = 511.65898556365914 ||V(L( $\beta$ , $\lambda$ ))||_2 = 258.88738725459575 gamma = 0.00078364164096 ACC = 0.98 PRE = 0.984126394045 REC = 0.984126394045
k = 8 t = 30.066906929 L( $\beta$ , $\lambda$ ) = 484.9422819846264 ||V(L( $\beta$ , $\lambda$ ))||_2 = 276.7354829899021 gamma = 0.00078364164096 ACC = 0.984 PRE = 0.976277372263 REC = 0.994423791822
k = 9 t = 33.3298799992 L( $\beta$ , $\lambda$ ) = 467.33336504481593 ||V(L( $\beta$ , $\lambda$ ))||_2 = 297.33568125873046 gamma = 0.00078364164096 ACC = 0.979 PRE = 0.981378026071 REC = 0.979553903346
k = 10 t = 36.7253370285 L( $\beta$ , $\lambda$ ) = 449.768322792532 ||V(L( $\beta$ , $\lambda$ ))||_2 = 319.5367352237993 gamma = 0.000470184284576 ACC = 0.984 PRE = 0.979779411765 REC = 0.990706319703
k = 11 t = 39.7053449154 L( $\beta$ , $\lambda$ ) = 425.92454425676246 ||V(L( $\beta$ , $\lambda$ ))||_2 = 166.72120148928928 gamma = 0.002176782336 ACC = 0.983 PRE = 0.988742964349 REC = 0.976539393935
k = 12 t = 43.0653970165 L( $\beta$ , $\lambda$ ) = 405.56100000000004 ||V(L( $\beta$ , $\lambda$ ))||_2 = 38.19925965616004 gamma = 0.000470184284576 ACC = 0.986 PRE = 0.981238273921 REC = 0.98444691822
k = 13 t = 46.6290397007 L( $\beta$ , $\lambda$ ) = 38.19925965616004 ||V(L( $\beta$ , $\lambda$ ))||_2 = 147.3928648463 gamma = 0.000470184284576 ACC = 0.977 PRE = 0.943582323485 REC = 0.98142659405
k = 14 t = 49.4615418911 L( $\beta$ , $\lambda$ ) = 364.10371139630433 ||V(L( $\beta$ , $\lambda$ ))||_2 = 357.7897348748522 gamma = 0.000470184284576 ACC = 0.989 PRE = 0.985267034991 REC = 0.994423791822
k = 15 t = 52.2741019726 L( $\beta$ , $\lambda$ ) = 335.5335812252476 ||V(L( $\beta$ , $\lambda$ ))||_2 = 119.37962057442935 gamma = 0.00362797056 ACC = 0.989 PRE = 0.994371482176 REC = 0.985130115524
k = 16 t = 55.691090107 L( $\beta$ , $\lambda$ ) = 321.40210325045234 ||V(L( $\beta$ , $\lambda$ ))||_2 = 359.39323294505124 gamma = 0.000470184284576 ACC = 0.994 PRE = 0.992592592593 REC = 0.996282527881
k = 17 t = 58.0605199337 L( $\beta$ , $\lambda$ ) = 290.83776835018057 ||V(L( $\beta$ , $\lambda$ ))||_2 = 93.27830407238 gamma = 0.01679616 ACC = 0.99 PRE = 0.994382022472 REC = 0.98698847584
k = 18 t = 61.3027141094 L( $\beta$ , $\lambda$ ) = 238.23489047134834 ||V(L( $\beta$ , $\lambda$ ))||_2 = 427.8373373957826 gamma = 0.00078364164096 ACC = 0.997 PRE = 0.99455713494 REC = 1.0
k = 19 t = 64.4026780128 L( $\beta$ , $\lambda$ ) = 183.51233512776878 ||V(L( $\beta$ , $\lambda$ ))||_2 = 104.68994981492128 gamma = 0.0013060694016 ACC = 0.997 PRE = 0.994454713494 REC = 1.0
Algorithm run for 20 iterations. Converged: False
Saving trained  $\beta$  in beta

```

$\lambda = 1.0$

```
[wang_zife@c0013 Assignment3]$ python LogisticRegression.py mushrooms/mushrooms.train --testdata mushrooms/mushrooms.test --lam 1.0 --max_iter 20
Reading training data from mushrooms/mushrooms.train
Read 7416 data points with 117 features in total
Reading test data from mushrooms/mushrooms.test
Read 1000 data points with 115 features
Training on data from mushrooms/mushrooms.train with : - 1.0 , - 0.1 , max iter = 20
k = 0 t = 3.2123111092 L(s..k) = 5148.379491033203 ||(L(s..k))||_2 = 4273.548233026341 gamma = 0.000470184984576 ACC = 0.909 PRE = 0.855325914149 REC = 1.0
k = 1 t = 3.2123111092 L(s..k) = 5148.379491033203 ||(L(s..k))||_2 = 4273.548233026341 gamma = 0.000470184984576 ACC = 0.921 PRE = 0.855325914149 REC = 0.968401486989
k = 2 t = 3.2123111092 L(s..k) = 5148.379491033203 ||(L(s..k))||_2 = 4273.548233026341 gamma = 0.000470184984576 ACC = 0.970 PRE = 0.965725765766 REC = 0.99628327881
k = 3 t = 12.9323890209 L(s..k) = 800.6863196497471 ||(L(s..k))||_2 = 1011.3003853022815 gamma = 0.000470184984576 ACC = 0.974 PRE = 0.981230307519 REC = 0.970260123048
k = 4 t = 16.285713007 L(s..k) = 674.8844560320292 ||(L(s..k))||_2 = 665.3000081264079 gamma = 0.000470184984576 ACC = 0.975 PRE = 0.97042201835 REC = 0.983271375465
k = 5 t = 19.5950909042 L(s..k) = 666.7689904402055 ||(L(s..k))||_2 = 407.56767081051424 gamma = 0.000470184984576 ACC = 0.98 PRE = 0.981412639405 REC = 0.981412639405
k = 6 t = 22.6482651238 L(s..k) = 566.9918768399507 ||(L(s..k))||_2 = 457.89658750961775 gamma = 0.000470184984576 ACC = 0.984 PRE = 0.976277372263 REC = 0.994423791822
k = 7 t = 25.9726719856 L(s..k) = 537.8265214670596 ||(L(s..k))||_2 = 432.0395827372417 gamma = 0.000470184984576 ACC = 0.98 PRE = 0.981412639405 REC = 0.981412639405
k = 8 t = 29.1943149567 L(s..k) = 498.52726352261834 ||(L(s..k))||_2 = 206.81000518374843 gamma = 0.000470184984576 ACC = 0.984 PRE = 0.976277372263 REC = 0.994423791822
k = 9 t = 32.609400876 L(s..k) = 480.771693495731 ||(L(s..k))||_2 = 347.81000518370663 gamma = 0.000470184984576 ACC = 0.984 PRE = 0.979784797048 REC = 0.98698847584
k = 10 t = 35.5935610283 L(s..k) = 462.817797675863 ||(L(s..k))||_2 = 359.167400861842 gamma = 0.000470184984576 ACC = 0.985 PRE = 0.976329582878 REC = 0.996282527881
k = 11 t = 39.008451023 L(s..k) = 404.8058260464825 ||(L(s..k))||_2 = 379.67900000000006 gamma = 0.000470184984576 ACC = 0.985 PRE = 0.981538379378 REC = 0.990706319703
k = 12 t = 42.4188509814 L(s..k) = 390.90000000000007 ||(L(s..k))||_2 = 379.67900000000006 gamma = 0.000470184984576 ACC = 0.985 PRE = 0.97205857878 REC = 0.996282527881
k = 13 t = 45.4102230072 L(s..k) = 390.90000000000007 ||(L(s..k))||_2 = 295.196532913116 gamma = 0.000470184984576 ACC = 0.989 PRE = 0.987084870849 REC = 0.994423791822
k = 14 t = 48.8861303333 L(s..k) = 369.9001644540126 ||(L(s..k))||_2 = 188.30978241611934 gamma = 0.0004661576 ACC = 0.987 PRE = 0.976406533575 REC = 1.0
k = 15 t = 51.3794751167 L(s..k) = 348.48549899718256 ||(L(s..k))||_2 = 390.5459898808085 gamma = 0.000470184984576 ACC = 0.992 PRE = 0.994402985975 REC = 0.990706319703
k = 16 t = 54.652310133 L(s..k) = 322.42985665680954 ||(L(s..k))||_2 = 256.7968251280171 gamma = 0.000470184984576 ACC = 0.994 PRE = 0.990747907749 REC = 0.998141263941
k = 17 t = 57.9140601158 L(s..k) = 301.3780633555753 ||(L(s..k))||_2 = 114.12659985625002 gamma = 0.0013060694016 ACC = 0.994 PRE = 0.992662846588 REC = 0.998141263941
k = 18 t = 60.9546010494 L(s..k) = 297.65146407090236 ||(L(s..k))||_2 = 154.96136506980739 gamma = 0.00078364164096 ACC = 0.995 PRE = 0.99443137291 REC = 0.996282527881
k = 19 t = 64.1429350376 L(s..k) = 297.65146407090236 ||(L(s..k))||_2 = 154.96136506980739 gamma = 0.00078364164096 ACC = 0.995 PRE = 0.99443137291 REC = 0.996282527881
Algorithm ran for 20 iterations. Converged: False
Saving trained s in beta
```

$\lambda = 10.0$

```
[wang_zife@c0013 Assignment3]$ python LogisticRegression.py mushrooms/mushrooms.train --testdata mushrooms/mushrooms.test --lam 10.0 --max_iter 20
Reading training data from mushrooms/mushrooms.train
Read 7416 data points with 117 features in total
Reading test data from mushrooms/mushrooms.test
Read 1000 data points with 115 features
Training on data from mushrooms/mushrooms.train with : - 10.0 , - 0.1 , max iter = 20
k = 0 t = 3.51563102837 L(s..k) = 5148.379491033203 ||(L(s..k))||_2 = 4273.548233026341 gamma = 0.000470184984576 ACC = 0.909 PRE = 0.855325914149 REC = 1.0
k = 1 t = 7.3973791351 L(s..k) = 2557.3691386544427 ||(L(s..k))||_2 = 3279.791419606979 gamma = 0.000169266594447 ACC = 0.921 PRE = 0.893653516295 REC = 0.968401486898
k = 2 t = 10.281635946 L(s..k) = 1580.467069310575 ||(L(s..k))||_2 = 786.7239775788153 gamma = 0.003627970956 ACC = 0.985 PRE = 0.965827338129 REC = 0.998141263941
k = 3 t = 13.9193339346 L(s..k) = 1064.089362549908 ||(L(s..k))||_2 = 1329.316239138266 gamma = 0.00028210996746 ACC = 0.976 PRE = 0.975925925926 REC = 0.975933903346
k = 4 t = 16.9486320019 L(s..k) = 826.628992958825 ||(L(s..k))||_2 = 322.7135610904411 gamma = 0.002176782336 ACC = 0.984 PRE = 0.97627372263 REC = 0.994423791822
k = 5 t = 20.4222066787 L(s..k) = 766.9113008389959 ||(L(s..k))||_2 = 649.8966106405724 gamma = 0.000470184984576 ACC = 0.977 PRE = 0.981384611215 REC = 0.975836431227
k = 6 t = 23.901219377 L(s..k) = 701.5728340244931 ||(L(s..k))||_2 = 299.23404558022776 gamma = 0.000470184984576 ACC = 0.982 PRE = 0.9779411767474 REC = 0.988447583643
k = 7 t = 27.001219377 L(s..k) = 687.53487624044 ||(L(s..k))||_2 = 174.5686166464094 gamma = 0.000470184984576 ACC = 0.983 PRE = 0.979774211765 REC = 0.98774565
k = 8 t = 30.703559578 L(s..k) = 677.5357426513 ||(L(s..k))||_2 = 198.724794745743 gamma = 0.000470184984576 ACC = 0.984 PRE = 0.980774211765 REC = 0.9870631703
k = 9 t = 33.893295262 L(s..k) = 658.8198570324075 ||(L(s..k))||_2 = 138.7221592157115 gamma = 0.0013060694015 ACC = 0.983 PRE = 0.981515711645 REC = 0.98698847584
k = 10 t = 37.362329062 L(s..k) = 659.5846528848683 ||(L(s..k))||_2 = 232.29008733610144 gamma = 0.000470184984576 ACC = 0.986 PRE = 0.979853479853 REC = 0.994423791822
k = 11 t = 40.7058670521 L(s..k) = 649.5221675986354 ||(L(s..k))||_2 = 130.53627196605 gamma = 0.00078364164096 PRE = 0.983394833948 REC = 0.990706319703
k = 12 t = 44.1735649109 L(s..k) = 644.27745474574904 ||(L(s..k))||_2 = 152.9467941217904 gamma = 0.00078364164096 PRE = 0.979853479853 REC = 0.994423791822
k = 13 t = 47.3565020561 L(s..k) = 638.7149242589148 ||(L(s..k))||_2 = 94.40178656592078 gamma = 0.0013060694016 ACC = 0.988 PRE = 0.987037037037 REC = 0.990706319703
k = 14 t = 50.8273594948 ||(L(s..k))||_2 = 633.8041949112851 ||(L(s..k))||_2 = 173.5059895809508 gamma = 0.00078364164096 PRE = 0.979853479853 REC = 0.994423791822
k = 15 t = 54.1522438526 L(s..k) = 627.9060050562742 ||(L(s..k))||_2 = 95.36541625530585 gamma = 0.00078364164096 ACC = 0.99 PRE = 0.987084870849 REC = 0.994423791822
k = 16 t = 57.4043043866 L(s..k) = 627.9060050562742 ||(L(s..k))||_2 = 101.1140695195644 gamma = 0.00078364164096 ACC = 0.984 PRE = 0.9816349918685 REC = 0.996282527881
k = 17 t = 60.962499943 L(s..k) = 621.8652531701674 ||(L(s..k))||_2 = 127.02344628818398 gamma = 0.000470184984576 ACC = 0.99 PRE = 0.987084870849 REC = 0.994423791822
k = 18 t = 64.5304046466 L(s..k) = 618.177992595795 ||(L(s..k))||_2 = 74.27459588799452 gamma = 0.0013060694018 ACC = 0.989 PRE = 0.985486258552 REC = 0.994423791822
k = 19 t = 67.609679169 L(s..k) = 615.2920528686961 ||(L(s..k))||_2 = 126.35522082323711 gamma = 0.000470184984576 ACC = 0.991 PRE = 0.988909426987
Algorithm ran for 20 iterations. Converged: False
Saving trained s in beta
```

3(d)

We definitely want a model with (a) high precision and low recall, so that even if we could not find out all the edible mushrooms, we would minimize the possibility to misclassify poisonous mushroom as edible, which is extremely dangerous.

## Problem 4

4(a)

```
def getAllFeaturesRDD(dataRDD):
    """ Get all the features present in grouped dataset dataRDD.

    The input is:
        - dataRDD containing pairs of the form (SparseVector(x),y).

    The return value is an list containing the union of all unique features present in sparse vectors inside dataRDD.
    """
    all_features = dataRDD.map(lambda (x,y): x)\n        .reduce(add)\n\n    return all_features.keys()

def totalLossRDD(dataRDD,beta,lam = 0.0):
    """ Given a sparse vector beta and a dataset in the form of RDD compute the regularized total logistic loss :

        L(\beta) = \sum_{(x,y) in data} l(\beta;x,y) + \lambda ||\beta||^2

    Inputs are:
        - dataRDD: an RDD containing pairs of the form (x,y), where x is a sparse vector and y is a binary value
        - beta: a sparse vector \beta
        - lam: the regularization parameter \lambda

    The return value is a float number represents the total loss
    """
    loss = dataRDD.map(lambda (x,y): logisticLoss(beta, x, y))\n        .reduce(add)\n    return loss + lam * beta.dot(beta)
```

```

def gradTotalLossRDD(dataRDD,beta,lam = 0.0):
    """ Given a sparse vector beta and a dataset in the form of RDD, compute the gradient of regularized total logistic loss :
         $\nabla L(\beta) = \sum_{(x,y) \in \text{data}} \nabla l(\beta; x, y) + 2\lambda \beta$ 
    Inputs are:
        - data: an RDD containing pairs of the form  $(x, y)$ , where  $x$  is a sparse vector and  $y$  is a binary value
        - beta: a sparse vector  $\beta$ 
        - lam: the regularization parameter  $\lambda$ 
    The return value is a float number represents the gradient of total loss
    """
    gradLoss = dataRDD.map(lambda (x, y): gradLogisticLoss(beta, x, y))\
        .reduce(add)
    return gradLoss + 2.0*lam*beta

def test(dataRDD,beta):
    """ Output the quantities necessary to compute the accuracy, precision, and recall of the prediction of labels in a dataset under a given  $\beta$ .
    The accuracy (ACC), precision (PRE), and recall (REC) are defined in terms of the following sets:
        P = datapoints  $(x,y)$  in data for which  $\langle \beta, x \rangle > 0$ 
        N = datapoints  $(x,y)$  in data for which  $\langle \beta, x \rangle \leq 0$ 
        TP = datapoints in  $(x,y)$  in P for which  $y=1$ 
        FP = datapoints in  $(x,y)$  in P for which  $y=-1$ 
        TN = datapoints in  $(x,y)$  in N for which  $y=-1$ 
        FN = datapoints in  $(x,y)$  in N for which  $y=1$ 
    For #XXX the number of elements in set XXX, the accuracy, precision, and recall of parameter vector  $\beta$  over data are defined as:
        ACC( $\beta$ ,data) = (#TP+#TN) / (#P + #N)
        PRE( $\beta$ ,data) = #TP / (#TP + #FP)
        REC( $\beta$ ,data) = #TP / (#TP + #FN)
    Inputs are:
        - dataRDD: an RDD containing pairs of the form  $(x,y)$ 
        - beta: vector  $\beta$ 
    The return values are
        - ACC, PRE, REC
    """
    P_set = dataRDD.filter(lambda (x,y): x.dot(beta) > 0).cache()
    N_set = dataRDD.filter(lambda (x,y): x.dot(beta) <= 0).cache()
    P = P_set.count()
    N = N_set.count()
    TP = P_set.filter(lambda (x,y): y == 1).count()
    FP = P - TP
    TN = N_set.filter(lambda (x,y): y == -1).count()
    FN = N - TN
    ACC = (TP + TN) / float(P + N)
    PRE = TP / float(TP + FP)
    REC = TP / float(TP + FN)
    return ACC, PRE, REC

def train(dataRDD,beta_0,lambda,max_iter,eps,test_data=None):
    """ Perform parallel logistic regression:
        to minimize
             $L(\beta) = \sum_{(x,y) \in \text{data}} l(\beta; x, y) + \lambda \|\beta\|_2^2$ 
    where the inputs are:
        - dataRDD: an rdd containing pairs of the form  $(x,y)$  as train data
        - beta_0: the starting vector  $\beta$ 
        - lambda: is the regularization parameter  $\lambda$ 
        - max_iter: maximum number of iterations of gradient descent
        - eps: upper bound on the L2 norm of the gradient
        - test_data: an rdd containing pairs of the form  $(x,y)$  as test data
    The function returns:
        - beta: the trained  $\beta$ ,
        - gradNorm: the norm of the gradient at the trained  $\beta$ , and
        - k: the number of iterations performed
    """
    k = 0
    gradNorm = 2*eps
    beta = beta_0
    start = time()
    while k<max_iter and gradNorm > eps:
        obj = totalLossRDD(dataRDD,beta,lambda)

        grad = gradTotalLossRDD(dataRDD,beta,lambda)
        gradNormSq = grad.dot(grad)
        gradNorm = np.sqrt(gradNormSq)

        fun = lambda x: totalLossRDD(dataRDD,x,lambda)
        gamma = lineSearch(fun,beta,grad,obj,gradNormSq)

        beta = beta - gamma * grad
        if test_data == None:
            print 'k = ',k,' \t t = ',time()-start,' \t L(\beta_k) = ',obj,' \t ||\nabla L(\beta_k)||_2 = ',gradNorm,' \t gamma = ',gamma
        else:
            acc,pre,rec = test(test_data,beta)
            print 'k = ',k,' \t t = ',time()-start,' \t L(\beta_k) = ',obj,' \t ||\nabla L(\beta_k)||_2 = ',gradNorm,' \t gamma = ',gamma,' \t ACC = ',acc,' \t PRE = ',pre,' \t REC = ',rec
        k = k + 1

    return beta,gradNorm,k

```

```

if __name__ == "__main__":
    parser = argparse.ArgumentParser(description = 'Parallel Logistic Regression.',formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    parser.add_argument('--traindata',default=None, help='Input file containing (x,y) pairs, used to train a logistic model')
    parser.add_argument('--testdata', default=None, help='Input file containing (x,y) pairs, used to test a logistic model')
    parser.add_argument('--lam', type=float, default=0.0, help='File where beta is stored (when training) and read from (when testing)')
    parser.add_argument('--max_iter', type=int, default=100, help='Maximum number of iterations')
    parser.add_argument('--eps', type=float, default=0.1, help='E-tolerance. If the l2_norm gradient is smaller than ε, gradient descent terminates.')

    args = parser.parse_args()

    sc = SparkContext(appName='Parallel Logistic Regression')

    print 'Reading training data from',args.traindata
    traindata = readDataRDD(args.traindata, sc).cache()
    print 'Read',traindata.count(),'data points with',len(getAllFeaturesRDD(traindata)), 'features in total'

    if args.testdata is not None:
        print 'Reading test data from',args.testdata
        testdata = readDataRDD(args.testdata, sc).cache()
        print 'Read',testdata.count(),'data points with',len(getAllFeaturesRDD(testdata)), 'features'
    else:
        testdata = None

    beta0 = SparseVector({})

    print 'Training on data from',args.traindata,'with λ =',args.lam,'ε =',args.eps,'max iter = ',args.max_iter
    beta, gradNorm, k = train(traindata,beta0,lam=args.lam,max_iter=args.max_iter,eps=args.eps,test_data=testdata)
    print 'Algorithm ran for:',k,'iterations. Converged:',gradNorm<args.eps
    print 'Saving trained β in',args.beta
    writeBeta(args.beta,beta)

```

## 4(b) Results on mushrooms dataset

This is the previous result using LogisticRegression.py

```

lwang_zifei@0013 Assignment315 python LogisticRegression.py mushrooms/mushrooms.train --testdata mushrooms/mushrooms.test --lam 0.0 --max_iter 20
Reading training data from mushrooms/mushrooms.train
Read 7416 data points with 117 features in total
Reading test data from mushrooms/mushrooms.test
Read 1000 data points with 115 features
Training on data from mushrooms/mushrooms.train with λ = 0.0 , ε = 0.1 , max iter = 20
k = 0 t = 3.35465294189 L(ρ,λ) = 5140.379491032203 ||V(ρ,λ)||_2 = 4273.548233026041 gamma = 0.000470184984576 ACC = 0.909 PRE = 0.855325914149 REC = 1.0
k = 1 t = 7.9175490497 L(ρ,λ) = 2516.993994492876 ||V(ρ,λ)||_2 = 978.986464992974 gamma = 0.000169266594447 ACC = 0.92 PRE = 0.892123827671 REC = 0.968401486989
k = 2 t = 10.0537118912 L(ρ,λ) = 1536.62968915103 ||V(ρ,λ)||_2 = 802.3372840311758 gamma = 0.00362797056 ACC = 0.979 PRE = 0.96576576576 REC = 0.996282572781
k = 3 t = 13.516380196 L(ρ,λ) = 772.3391480787859 ||V(ρ,λ)||_2 = 302.5497291421832 gamma = 0.00070184984576 ACC = 0.974 PRE = 0.981293007519 REC = 0.970260223048
k = 4 t = 16.908779199 L(ρ,λ) = 648.968995239955 ||V(ρ,λ)||_2 = 636.7301783321842 gamma = 0.000470184984576 ACC = 0.976 PRE = 0.972426470588 REC = 0.983271375465
k = 5 t = 20.057470315 L(ρ,λ) = 584.8639856207043 ||V(ρ,λ)||_2 = 398.50093415046 gamma = 0.000470184984576 ACC = 0.975 PRE = 0.972426470588 REC = 0.983271375465
k = 6 t = 23.111111111 L(ρ,λ) = 554.8639856207043 ||V(ρ,λ)||_2 = 274.3805736703113 gamma = 0.000470184984576 ACC = 0.971 PRE = 0.97790552466 REC = 0.9698848875384
k = 7 t = 26.8184449009 L(ρ,λ) = 511.6589856536591 ||V(ρ,λ)||_2 = 258.8873872549575 gamma = 0.00073343164096 ACC = 0.98 PRE = 0.981412639495 REC = 0.981412639495
k = 8 t = 30.0669463929 L(ρ,λ) = 468.8639856207043 ||V(ρ,λ)||_2 = 276.735828399021 gamma = 0.00073343164096 ACC = 0.984 PRE = 0.97627372263 REC = 0.994423791822
k = 9 t = 33.3298799992 L(ρ,λ) = 467.3333650481593 ||V(ρ,λ)||_2 = 297.33568125873046 gamma = 0.00073343164096 ACC = 0.979 PRE = 0.981378026071 REC = 0.979553903346
k = 10 t = 36.72523702852 L(ρ,λ) = 319.5307352273973 ||V(ρ,λ)||_2 = 139.5307352273973 gamma = 0.00073343164096 ACC = 0.984 PRE = 0.979779411765 REC = 0.99076319703
k = 11 t = 39.7053449154 L(ρ,λ) = 425.9245442967624 ||V(ρ,λ)||_2 = 166.7220148298298 gamma = 0.002176782336 ACC = 0.983 PRE = 0.988742964353 REC = 0.979553903346
k = 12 t = 43.008263970165 L(ρ,λ) = 405.9943038552676 ||V(ρ,λ)||_2 = 104.9943038552676 gamma = 0.000270184984576 ACC = 0.985 PRE = 0.97853479853 REC = 0.994423791822
k = 13 t = 46.0599400997 L(ρ,λ) = 374.25996474179607 ||V(ρ,λ)||_2 = 147.1673751224483 gamma = 0.002176782336 ACC = 0.987 PRE = 0.99435022486 REC = 0.981412639495
k = 14 t = 49.461518911 L(ρ,λ) = 364.1037113639485 ||V(ρ,λ)||_2 = 74.3805736703113 gamma = 0.000270184984576 ACC = 0.989 PRE = 0.985267034991 REC = 0.994423791822
k = 15 t = 52.2741019726 L(ρ,λ) = 359.533581252476 ||V(ρ,λ)||_2 = 119.3765075442935 gamma = 0.00362797056 ACC = 0.989 PRE = 0.99437482176 REC = 0.98513011524
k = 16 t = 55.3027141094 L(ρ,λ) = 321.4073382407238 ||V(ρ,λ)||_2 = 389.3027141094 gamma = 0.000270184984576 ACC = 0.99 PRE = 0.994382622472 REC = 0.996282527881
k = 17 t = 58.06651190337 L(ρ,λ) = 298.8373833818857 ||V(ρ,λ)||_2 = 93.278382407238 gamma = 0.01679616 ACC = 0.99 PRE = 0.994382622472 REC = 0.9698847384
k = 18 t = 61.3027141094 L(ρ,λ) = 238.23480047134834 ||V(ρ,λ)||_2 = 427.8373833818857 gamma = 0.00073343164096 ACC = 0.997 PRE = 0.99445713494 REC = 1.0
k = 19 t = 64.4026780218 L(ρ,λ) = 183.51233512776878 ||V(ρ,λ)||_2 = 184.68994981492128 gamma = 0.0013060694016 ACC = 0.997 PRE = 0.99445713494 REC = 1.0
Algorithm ran for 20 iterations. Converged: False
Saving trained a in beta

```

Here below is the result using parallel implementation

```

lwang_zifei@0010 Assignment315 python ParallelLogisticRegression.py mushrooms/mushrooms.train --testdata mushrooms/mushrooms.test --lam 0.0 --max_iter 20
2019-03-18 00:08:02 WARN NativeCodeLoader:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Reading training level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Reading training data from mushrooms/mushrooms.train
Read 7416 data points with 117 features in total
Reading test data from mushrooms/mushrooms.train
Read 1000 data points with 115 features
Training on data from mushrooms/mushrooms.train with λ = 0.0 , ε = 0.1 , max iter = 20
k = 0 t = 1.60661108539 L(ρ,λ) = 5140.379491032573 ||V(ρ,λ)||_2 = 4273.548233026041 gamma = 0.000470184984576 ACC = 0.909 PRE = 0.855325914149 REC = 1.0
k = 1 t = 7.9175490497 L(ρ,λ) = 2516.993994492844 ||V(ρ,λ)||_2 = 978.986464992974 gamma = 0.000169266594447 ACC = 0.92 PRE = 0.892123827671 REC = 0.968401486989
k = 2 t = 10.8794829845 L(ρ,λ) = 1536.62968915103 ||V(ρ,λ)||_2 = 802.3372840311758 gamma = 0.00362797056 ACC = 0.979 PRE = 0.96576576576 REC = 0.996282572781
k = 3 t = 14.473999234 L(ρ,λ) = 772.3391480787859 ||V(ρ,λ)||_2 = 988.5492791241635 gamma = 0.00070184984576 ACC = 0.974 PRE = 0.981293007519 REC = 0.970260223048
k = 4 t = 18.3914749622 L(ρ,λ) = 648.968995239362 ||V(ρ,λ)||_2 = 636.7301783321427 gamma = 0.000470184984576 ACC = 0.976 PRE = 0.972426470588 REC = 0.983271375465
k = 5 t = 21.7866768837 L(ρ,λ) = 583.129303296539 ||V(ρ,λ)||_2 = 390.12584418729 gamma = 0.000270184984576 ACC = 0.975 PRE = 0.981238273921 REC = 0.972118959108
k = 6 t = 25.2595560551 L(ρ,λ) = 405.9943038552676 ||V(ρ,λ)||_2 = 474.38057367033964 gamma = 0.000270184984576 ACC = 0.981 PRE = 0.97990552486 REC = 0.98698847584
k = 7 t = 28.584913969 L(ρ,λ) = 351.6589856536261 ||V(ρ,λ)||_2 = 258.8873872549576 gamma = 0.00073343164096 ACC = 0.98 PRE = 0.98142639405 REC = 0.98142639405
k = 8 t = 32.0210180239 L(ρ,λ) = 464.942281946284 ||V(ρ,λ)||_2 = 276.7354829899119 gamma = 0.00073343164096 ACC = 0.984 PRE = 0.97627732263 REC = 0.994423791822
k = 9 t = 35.5006780624 L(ρ,λ) = 467.3333650481593 ||V(ρ,λ)||_2 = 297.33568125873046 gamma = 0.00073343164096 ACC = 0.997 PRE = 0.991308726071 REC = 0.979553903346
k = 10 t = 39.1885109365 L(ρ,λ) = 449.7683272253473 ||V(ρ,λ)||_2 = 310.397352273975 gamma = 0.000270184984576 ACC = 0.984 PRE = 0.97779411765 REC = 0.979553903346
k = 11 t = 42.1263969125 L(ρ,λ) = 405.9943038552676 ||V(ρ,λ)||_2 = 166.7220148298298 gamma = 0.002176782336 ACC = 0.983 PRE = 0.988742964353 REC = 0.979553903346
k = 12 t = 45.8031110733 L(ρ,λ) = 405.9943038552711 ||V(ρ,λ)||_2 = 388.199259061811 gamma = 0.000270184984576 ACC = 0.986 PRE = 0.979853479853 REC = 0.994423791822
k = 13 t = 49.0535988928 L(ρ,λ) = 374.2599647417946 ||V(ρ,λ)||_2 = 147.1673751224483 gamma = 0.002176782336 ACC = 0.987 PRE = 0.99435022486 REC = 0.981412639405
k = 14 t = 52.712299671 L(ρ,λ) = 364.1037113639485 ||V(ρ,λ)||_2 = 357.7973487486793 gamma = 0.000270184984576 ACC = 0.989 PRE = 0.994371482176 REC = 0.98513011524
k = 15 t = 55.5561490059 L(ρ,λ) = 353.533581252484 ||V(ρ,λ)||_2 = 119.37562304743033 gamma = 0.00362797056 ACC = 0.989 PRE = 0.994371482176 REC = 0.98513011524
k = 16 t = 58.8773899078 L(ρ,λ) = 321.4073382407238 ||V(ρ,λ)||_2 = 359.3932835496896 gamma = 0.00073343164096 ACC = 0.994 PRE = 0.992592592593 REC = 0.996282527881
k = 17 t = 61.316818951 L(ρ,λ) = 290.8377683501805 ||V(ρ,λ)||_2 = 93.278382407238 gamma = 0.01679616 ACC = 0.99 PRE = 0.994382622472 REC = 0.98698847584
k = 18 t = 64.6839189529 L(ρ,λ) = 238.23480047135536 ||V(ρ,λ)||_2 = 427.83733795149 gamma = 0.00073343164096 ACC = 0.997 PRE = 0.99445713494 REC = 1.0
k = 19 t = 67.878740421 L(ρ,λ) = 183.51233512776867 ||V(ρ,λ)||_2 = 104.6899498149294 gamma = 0.0013060694016 ACC = 0.997 PRE = 0.99445713494 REC = 1.0
Algorithm ran for 20 iterations. Converged: False
Saving trained a in beta

```

Here the LogisticRegression.py is a little bit faster (64.4 seconds V.S. 67.8 seconds), but I think the difference is not that outstanding. So the running times are comparable, maybe if I optimize the parallel version better, I could get better results. But since the dimensionality of the data is not that big, a simple for loop may outperform map-reduce procedure by not having communication and data transfer time etc.

#### 4(c) Results on newsgroups dataset

This is the previous result using LogisticRegression.py

```
[wang_zifeic0010 Assignment3] python LogisticRegression.py newsgroups/news.train --testdata newsgroups/news.test --l1m 0.0 --max_iter 20
Reading training data from newsgroups/news.train
Read 1191 data points with 13361 features in total
Reading test data from newsgroups/news.test
Read 793 data points with 10824 features
Training on data from newsgroups.news.train with: <= 0.0 , <= 0.1 , max iter = 20
k = 0 t = 8.37953710556 L(ρ, k) = 825.5382920468972 |||L(ρ, k)||_2 = 583.679402529032 gamma = 0.008246176 ACC = 0.86380822825 PRE = 0.788 REC = 0.99449494949
k = 1 t = 1.9015000000000002 L(ρ, k) = 235.9977706984395 |||L(ρ, k)||_2 = 232.5729965119225 gamma = 0.00362797956 ACC = 0.923076923075 PRE = 0.98833819242 REC = 0.856060606061
k = 2 t = 0.291999054 L(ρ, k) = 165.31348191359206 |||L(ρ, k)||_2 = 170.445780558026 gamma = 0.00362797956 ACC = 0.95838576419 PRE = 0.934967312349 REC = 0.979797979798
k = 3 t = 0.381224759519 L(ρ, k) = 127.0884866361496 |||L(ρ, k)||_2 = 79.13053388943794 gamma = 0.0063046176 ACC = 0.956646910467 PRE = 0.9715025906706 REC = 0.946969696967
k = 4 t = 46.6291936076 L(ρ, k) = 113.2166966999398 |||L(ρ, k)||_2 = 58.86316891572183 gamma = 0.010677696 ACC = 0.954692774275 PRE = 0.928571428571 REC = 0.984848484848
k = 5 t = 55.0932440758 L(ρ, k) = 102.08000381733448 |||L(ρ, k)||_2 = 73.28323405204426 gamma = 0.0063046176 ACC = 0.965952080706 PRE = 0.971867007673 REC = 0.959595959596
k = 6 t = 63.4027272171 L(ρ, k) = 88.403489222394 |||L(ρ, k)||_2 = 39.94201967449916 gamma = 0.01679616 ACC = 0.956646910467 PRE = 0.93961352657 PRE = 0.98232323232
k = 7 t = 71.7419559956 L(ρ, k) = 79.66467843478928 |||L(ρ, k)||_2 = 58.51794902886756 gamma = 0.016777696 ACC = 0.962168978562 PRE = 0.98670212766 REC = 0.936868686869
k = 8 t = 80.151540041 L(ρ, k) = 69.2890442476854 |||L(ρ, k)||_2 = 42.23288682113263 gamma = 0.010677696 ACC = 0.9709986216898 PRE = 0.965087281796 REC = 0.977272727273
k = 9 t = 88.4505119987 L(ρ, k) = 61.501004619181272 gamma = 0.0279536
Algorithm ran for 20 iterations. Converged: False
Saving trained a in beta
```

Here below is the result using parallel implementation

```
[wang_zifeic0010 Assignment3] python ParallelLogisticRegression.py newsgroups/news.train --testdata newsgroups/news.test --l1m 0.0 --max_iter 20
2019-03-18 00:23:56 WARN NativeCodeCoder:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Reading training data from newsgroups/news.train
Read 1191 data points with 13361 features in total
Reading test data from newsgroups/news.test
Read 793 data points with 10824 features
Training on data from newsgroups/train with <= 0.0 , <= 0.1 , max iter = 20
k = 0 t = 6.86738610268 L(ρ, k) = 825.5382920468972 |||L(ρ, k)||_2 = 583.679402529032 gamma = 0.006466176 ACC = 0.862547288777 PRE = 0.786427145709 REC = 0.99449494949
k = 1 t = 14.3480689526 L(ρ, k) = 235.997706984395 |||L(ρ, k)||_2 = 298.527996018906 gamma = 0.00362797956 ACC = 0.923076923075 PRE = 0.98833819242 REC = 0.856060606061
k = 2 t = 22.108150000000002 L(ρ, k) = 165.31348191359197 |||L(ρ, k)||_2 = 170.445780558026 gamma = 0.00362797956 ACC = 0.95838576419 PRE = 0.934967312349 REC = 0.979797979798
k = 3 t = 29.189320874 L(ρ, k) = 127.0884866361497 |||L(ρ, k)||_2 = 79.13053388943804 gamma = 0.0063046176 ACC = 0.956646910467 PRE = 0.9715025906706 REC = 0.946969696967
k = 4 t = 36.230973959 L(ρ, k) = 113.2166966999398 |||L(ρ, k)||_2 = 58.86316891572172 gamma = 0.010677696 ACC = 0.954692774275 PRE = 0.928571428571 REC = 0.984848484848
k = 5 t = 43.4233418558 L(ρ, k) = 102.0800038173344 |||L(ρ, k)||_2 = 73.27832340520459 gamma = 0.006466176 ACC = 0.965952080706 PRE = 0.971867007673 REC = 0.99449494949
k = 6 t = 50.4233418557 L(ρ, k) = 88.403489222394 |||L(ρ, k)||_2 = 50.59166666666666 gamma = 0.010677696 ACC = 0.9709986216898 PRE = 0.965087281796 REC = 0.977272727273
k = 7 t = 56.1311939278 L(ρ, k) = 79.66467843478928 |||L(ρ, k)||_2 = 48.51794902886756 gamma = 0.01679616 ACC = 0.962168978562 PRE = 0.98670212766 REC = 0.936868686869
k = 8 t = 64.545298939 L(ρ, k) = 69.2890442476852 |||L(ρ, k)||_2 = 42.23288682113271 gamma = 0.010677696 ACC = 0.9709986216898 PRE = 0.965087281796 REC = 0.977272727273
k = 9 t = 71.0227710438 L(ρ, k) = 61.501004619181272 gamma = 0.0279536
Algorithm ran for 20 iterations. Converged: False
Saving trained a in beta
```

Here ParallelLogisticRegression.py is faster (132.2 seconds V.S. 170.3 seconds). The difference is considerable and it is because of the high dimensionality of the data make map-reduce more efficient than a single for loop.

## Problem 5

For convenience, I modified train() in both LogisticRegression.py and ParallelLogisticRegression.py to save certain variables for plotting. But I did not reflect the changes in the previous questions, so please refer to the final version attached.

Also plots are created on local machine using pyplot and Jupyter Notebook, an example snippet is below (all the plots follows this example except the parameters):

```
import matplotlib.pyplot as plt
import matplotlib
import numpy as np
import pickle
import os
root_path = '/Users/wangzifeng/Desktop/Courses@NEU/Parallel Programming/Homework/Ass3/'
time_lr = pickle.load(open(os.path.join(root_path, 'LR', 'time.pkl'), 'rb'))
time_plr = pickle.load(open(os.path.join(root_path, 'PLR', 'time.pkl'), 'rb'))
gradNorm_lr = pickle.load(open(os.path.join(root_path, 'LR', 'gradNorm.pkl'), 'rb'))
gradNorm_plr = pickle.load(open(os.path.join(root_path, 'PLR', 'gradNorm.pkl'), 'rb'))
acc_lr = pickle.load(open(os.path.join(root_path, 'LR', 'acc.pkl'), 'rb'))
acc_plr = pickle.load(open(os.path.join(root_path, 'PLR', 'acc.pkl'), 'rb'))
pre_lr = pickle.load(open(os.path.join(root_path, 'LR', 'pre.pkl'), 'rb'))
pre_plr = pickle.load(open(os.path.join(root_path, 'PLR', 'pre.pkl'), 'rb'))
rec_lr = pickle.load(open(os.path.join(root_path, 'LR', 'rec.pkl'), 'rb'))
rec_plr = pickle.load(open(os.path.join(root_path, 'PLR', 'rec.pkl'), 'rb'))

f = plt.figure(figsize=(12,7.2))

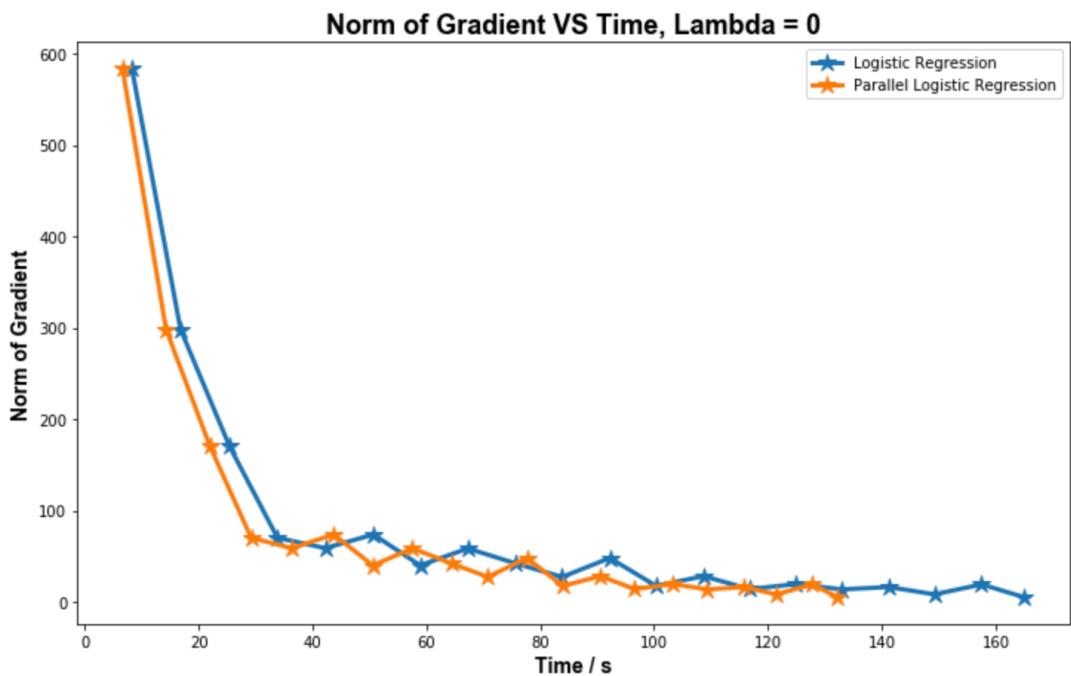
plt.plot(time_lr, gradNorm_lr, '-*', label = 'Logistic Regression', linewidth=3.0, markersize=12)
plt.plot(time_plr, gradNorm_plr, '-*', label = 'Parallel Logistic Regression', linewidth=3.0, markersize=12)

plt.legend()
plt.title("Norm of Gradient VS Time, Lambda = 0", fontsize='18', weight='bold', fontname='Arial')
plt.xlabel("Time / s", fontsize='14', weight='bold', fontname='Arial')
plt.ylabel("Norm of Gradient", fontsize='14', weight='bold', fontname='Arial')

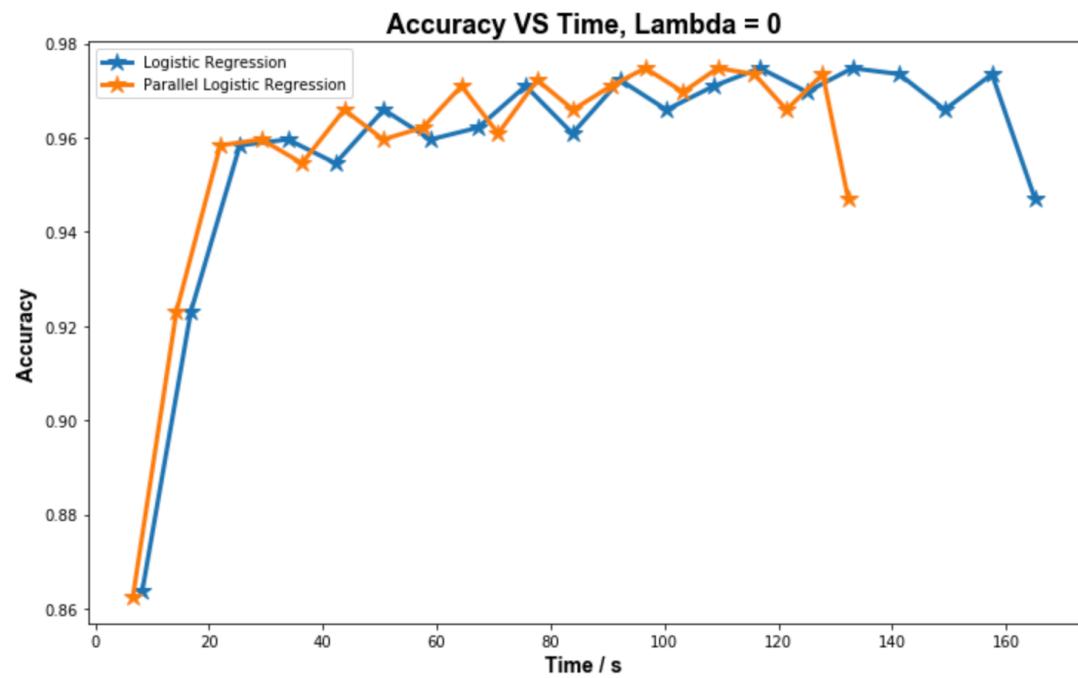
plt.show()
```

5(a) For each experiments, I run 20 iterations

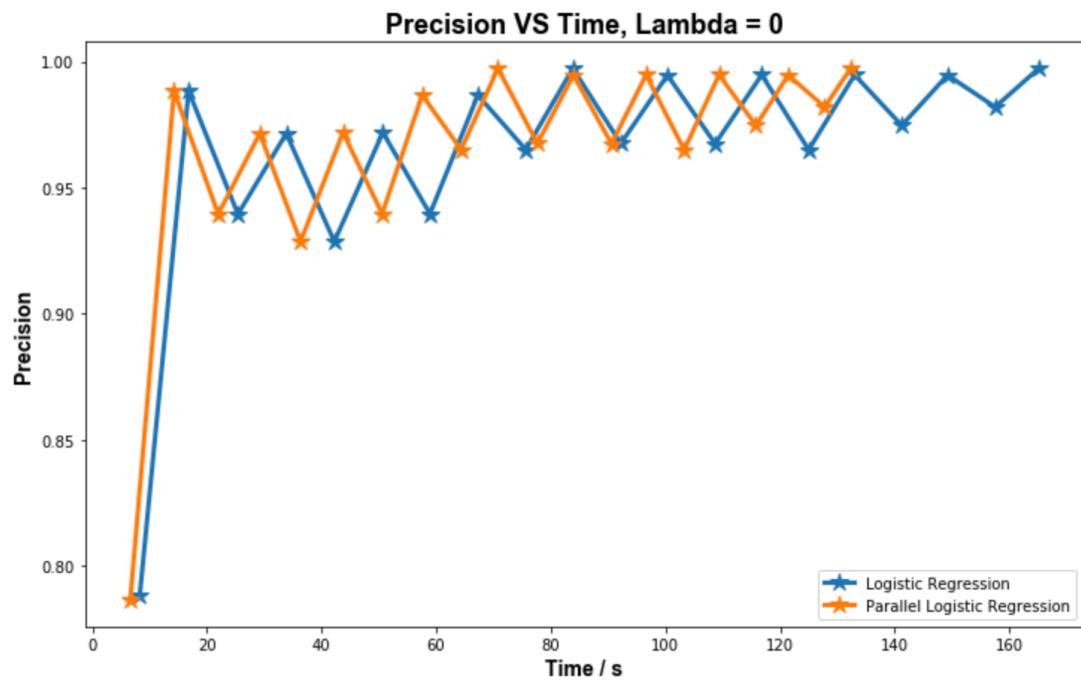
Plot for norm of gradient vs time.



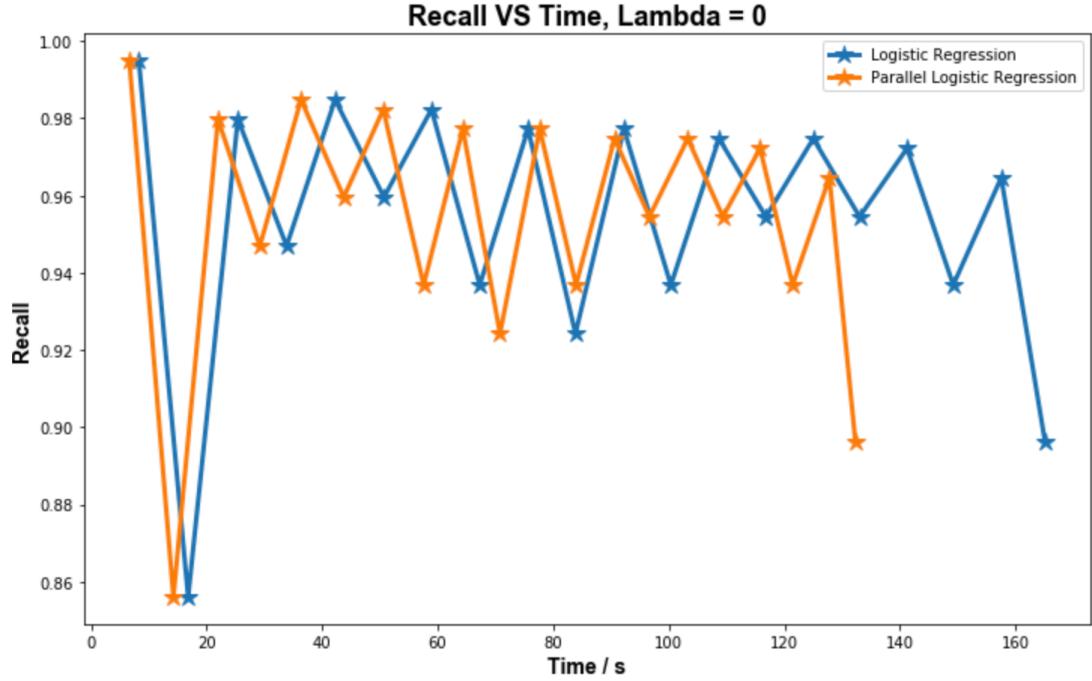
Plot for accuracy vs time



Plot for precision vs time

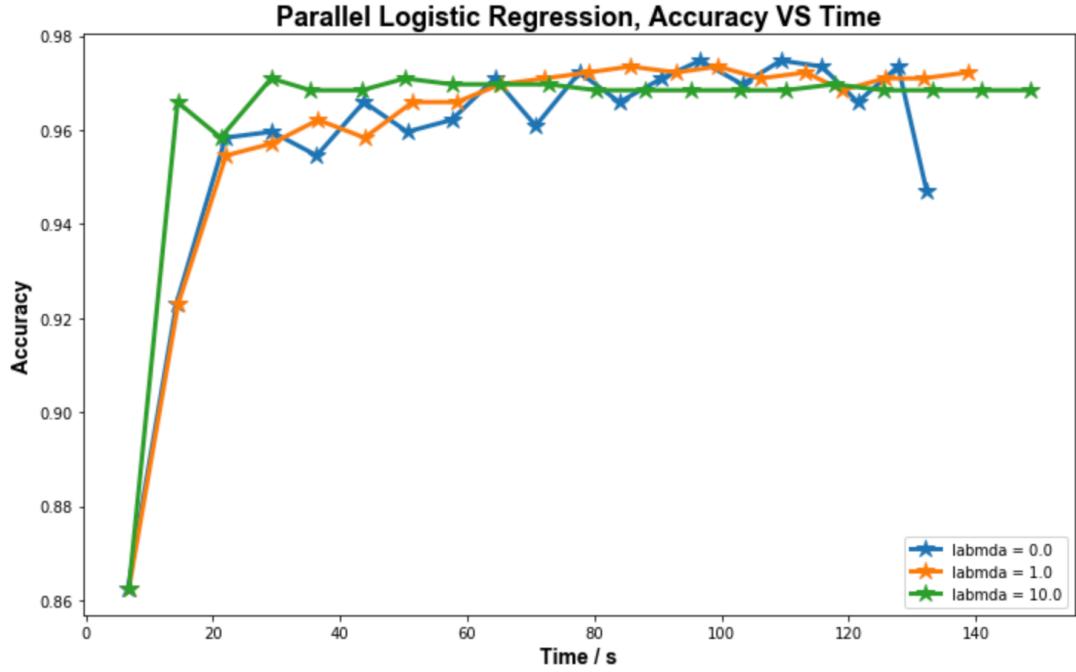


Plot for recall vs time



5(b) We choose  $\lambda$  from  $\{0.0, 1.0, 10.0\}$ , also for 20 epochs

For ParallelLogisticRegression.py



5(c)

From the plot in 5(b), we could see that the maximum accuracy is got when

$\lambda = 0.0$ , but still we need to verify which iteration is the best one (hard to tell from the plot), so I give the exact accuracy under each  $\lambda$  to find out which is the best.



As a result,  $\lambda = 0.0$ ,  $k = 13$  or  $\lambda = 0.0$   $k = 15$  seems to be the best combination, they give exactly the same performance.

Here I just choose  $\lambda = 0.0$ ,  $k = 13$  and find the corresponding beta below:

Features with 10 Most Positive Values

Feature	Value
doctor	0.739256993046
medic	0.654923809228
inform	0.64845288521
diseas	0.627412769331
treatment	0.594065145151
effect	0.588228768705
problem	0.542908430112
gordon	0.522539524406
peopl	0.509126766275
bank	0.505400995555

Features with 10 Most Negative Values

Feature	Value
basebal	-1.59643808354
game	-1.35316118574
team	-1.1894061731
player	-1.16455349865
win	-0.88657443609
plai	-0.885072905804
pitch	-0.79174957004
run	-0.780232214032
fan	-0.745623085766
yanke	-0.733704446503